

__data__modelling

July 10, 2023

```
[ ]: import warnings
warnings.filterwarnings('ignore')

# visuals
import tensorflow as tf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns
plt.style.use('seaborn')
%matplotlib inline

# models
import tensorflow as tf
from tensorflow import keras
from keras.wrappers.scikit_learn import KerasClassifier
from tensorflow.keras import Sequential, layers, callbacks
from tensorflow.keras.layers import Dense, LSTM, Dropout, GRU, Bidirectional

# validation
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import MinMaxScaler, StandardScaler

# metrics & evaluation
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.metrics import
    ↪mean_squared_error, mean_absolute_error, explained_variance_score, r2_score,
    ↪mean_absolute_percentage_error
from scipy import stats

print('Imports Complete')
```

Imports Complete

Using TensorFlow backend.

Data Exploration

This dataset has already been cleaned and preprocessed for modelling. This was done in the Data Preparation Notebook.

```
[ ]: # fix random seed for reproducibility
tf.random.set_seed(42)
```

```
[ ]: # load the dataset
df = pd.read_csv('mod_dec21_dec22_prepared_.csv')
# df = dataframe.values
df.head()
```

```
[ ]:
```

		timestamp	hour	sensor_id	P1	P2	pressure	temperature \
0	2021-12-01	11:13:48	11	67959.0	6.80	1.0	NaN	NaN
1	2021-12-01	11:16:16	11	67959.0	6.15	0.8	NaN	NaN
2	2021-12-01	11:18:45	11	67959.0	5.97	0.8	NaN	NaN
3	2021-12-01	11:23:44	11	67959.0	6.20	1.3	NaN	NaN
4	2021-12-01	11:26:15	11	67959.0	4.65	0.8	NaN	NaN

	humidity	day	month	year
0	NaN	1.0	12.0	2021.0
1	NaN	1.0	12.0	2021.0
2	NaN	1.0	12.0	2021.0
3	NaN	1.0	12.0	2021.0
4	NaN	1.0	12.0	2021.0

```
[ ]: df.tail()
```

```
[ ]:
```

		timestamp	hour	sensor_id	P1	P2	pressure \
1316743	2022-12-06	23:58:29	23	67961.5	19.250	10.93	101580.81
1316744	2022-12-06	23:58:43	23	67955.5	18.270	8.65	101676.44
1316745	2022-12-06	23:59:29	23	67959.0	33.670	15.60	101576.00
1316746	2022-12-06	23:59:30	23	67960.0	23.485	11.50	101475.56
1316747	2022-12-06	23:59:49	23	67993.5	13.300	7.40	101658.63

	temperature	humidity	day	month	year
1316743	3.950	100.00	6.0	12.0	2022.0
1316744	2.750	100.00	6.0	12.0	2022.0
1316745	2.835	100.00	6.0	12.0	2022.0
1316746	2.920	100.00	6.0	12.0	2022.0
1316747	7.510	74.25	6.0	12.0	2022.0

```
[ ]: df.shape
```

```
[ ]: (1316748, 11)
```

```
[ ]: # df["timestamp"] = pd.to_datetime(df[['year', 'month', 'day', 'hour']])

df['timestamp'] = pd.to_datetime(df['timestamp'],infer_datetime_format=True)

df.sort_values(by=['timestamp'], inplace=False)
```

```
[ ]:
```

	timestamp	hour	sensor_id	P1	P2	pressure \
0	2021-12-01 11:13:48	11	67959.0	6.800	1.00	NaN
1	2021-12-01 11:16:16	11	67959.0	6.150	0.80	NaN
2	2021-12-01 11:18:45	11	67959.0	5.970	0.80	NaN
3	2021-12-01 11:23:44	11	67959.0	6.200	1.30	NaN
4	2021-12-01 11:26:15	11	67959.0	4.650	0.80	NaN
...
1316743	2022-12-06 23:58:29	23	67961.5	19.250	10.93	101580.81
1316744	2022-12-06 23:58:43	23	67955.5	18.270	8.65	101676.44
1316745	2022-12-06 23:59:29	23	67959.0	33.670	15.60	101576.00
1316746	2022-12-06 23:59:30	23	67960.0	23.485	11.50	101475.56
1316747	2022-12-06 23:59:49	23	67993.5	13.300	7.40	101658.63

	temperature	humidity	day	month	year
0	NaN	NaN	1.0	12.0	2021.0
1	NaN	NaN	1.0	12.0	2021.0
2	NaN	NaN	1.0	12.0	2021.0
3	NaN	NaN	1.0	12.0	2021.0
4	NaN	NaN	1.0	12.0	2021.0
...
1316743	3.950	100.00	6.0	12.0	2022.0
1316744	2.750	100.00	6.0	12.0	2022.0
1316745	2.835	100.00	6.0	12.0	2022.0
1316746	2.920	100.00	6.0	12.0	2022.0
1316747	7.510	74.25	6.0	12.0	2022.0

[1316748 rows x 11 columns]

```
[ ]: df.head()
```

```
[ ]:
```

	timestamp	hour	sensor_id	P1	P2	pressure	temperature \
0	2021-12-01 11:13:48	11	67959.0	6.80	1.0	NaN	NaN
1	2021-12-01 11:16:16	11	67959.0	6.15	0.8	NaN	NaN
2	2021-12-01 11:18:45	11	67959.0	5.97	0.8	NaN	NaN
3	2021-12-01 11:23:44	11	67959.0	6.20	1.3	NaN	NaN
4	2021-12-01 11:26:15	11	67959.0	4.65	0.8	NaN	NaN

	humidity	day	month	year
0	NaN	1.0	12.0	2021.0
1	NaN	1.0	12.0	2021.0
2	NaN	1.0	12.0	2021.0

```
3      NaN  1.0   12.0  2021.0
4      NaN  1.0   12.0  2021.0
```

```
[ ]: df.tail()
```

```
[ ]:
      timestamp  hour  sensor_id    P1    P2  pressure \
1316743 2022-12-06 23:58:29    23    67961.5  19.250  10.93  101580.81
1316744 2022-12-06 23:58:43    23    67955.5  18.270   8.65  101676.44
1316745 2022-12-06 23:59:29    23    67959.0  33.670  15.60  101576.00
1316746 2022-12-06 23:59:30    23    67960.0  23.485  11.50  101475.56
1316747 2022-12-06 23:59:49    23    67993.5  13.300   7.40  101658.63

      temperature  humidity  day  month  year
1316743         3.950    100.00  6.0   12.0  2022.0
1316744         2.750    100.00  6.0   12.0  2022.0
1316745         2.835    100.00  6.0   12.0  2022.0
1316746         2.920    100.00  6.0   12.0  2022.0
1316747         7.510     74.25  6.0   12.0  2022.0
```

Overall, the data set has over a million records of measurements per minute with 11 columns.

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1316748 entries, 0 to 1316747
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   timestamp      1316748 non-null  datetime64[ns]
1   hour           1316748 non-null  int64
2   sensor_id      1316748 non-null  float64
3   P1             1316748 non-null  float64
4   P2             1316748 non-null  float64
5   pressure       1316255 non-null  float64
6   temperature    1316255 non-null  float64
7   humidity       1316255 non-null  float64
8   day            1316748 non-null  float64
9   month          1316748 non-null  float64
10  year           1316748 non-null  float64
dtypes: datetime64[ns](1), float64(9), int64(1)
memory usage: 110.5 MB
```

Below we set the group the data by hour. By doing that, we shrink the data to just over 8K rows.

```
[ ]: # print info to check conversion
df = df.set_index('timestamp').resample('H').mean() # set date as index or
      ↪ rest_index().resample('5min').mean()
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8893 entries, 2021-12-01 11:00:00 to 2022-12-06 23:00:00
Freq: H
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   hour             8330 non-null   float64
1   sensor_id        8330 non-null   float64
2   P1               8330 non-null   float64
3   P2               8330 non-null   float64
4   pressure         8309 non-null   float64
5   temperature      8309 non-null   float64
6   humidity         8309 non-null   float64
7   day              8330 non-null   float64
8   month            8330 non-null   float64
9   year             8330 non-null   float64
dtypes: float64(10)
memory usage: 764.2 KB
```

```
[ ]: df.head()
```

```
[ ]:
          hour  sensor_id      P1      P2  pressure \
timestamp
2021-12-01 11:00:00  11.0    67959.0    5.918750  0.921250      NaN
2021-12-01 12:00:00  12.0    67959.0    6.340000  0.891176      NaN
2021-12-01 13:00:00  13.0    67959.0    7.080000  0.987000      NaN
2021-12-01 14:00:00  14.0    67959.0   11.604091  1.667273      NaN
2021-12-01 15:00:00  15.0    67959.0   13.948500  2.034500      NaN

          temperature  humidity  day  month  year
timestamp
2021-12-01 11:00:00      NaN     NaN  1.0   12.0  2021.0
2021-12-01 12:00:00      NaN     NaN  1.0   12.0  2021.0
2021-12-01 13:00:00      NaN     NaN  1.0   12.0  2021.0
2021-12-01 14:00:00      NaN     NaN  1.0   12.0  2021.0
2021-12-01 15:00:00      NaN     NaN  1.0   12.0  2021.0
```

```
[ ]: df.tail()
```

```
[ ]:
          hour  sensor_id      P1      P2  pressure \
timestamp
2022-12-06 19:00:00  19.0  67966.827114    9.824515  5.174757  101612.213190
2022-12-06 20:00:00  20.0  67967.959732   12.462634  6.818859  101621.297785
2022-12-06 21:00:00  21.0  67968.278523   14.505705  7.973121  101606.998859
2022-12-06 22:00:00  22.0  67967.442029   16.657174  8.854710  101611.657500
```


```
2022-12-06 23:00:00    23.0  67967.000000  18.115521   9.781424  101604.348264
```

	temperature	humidity	day	month	year
timestamp					
2022-12-06 19:00:00	5.828451	92.733433	6.0	12.0	2022.0
2022-12-06 20:00:00	5.795403	92.526711	6.0	12.0	2022.0
2022-12-06 21:00:00	5.779866	92.406711	6.0	12.0	2022.0
2022-12-06 22:00:00	5.242500	92.941957	6.0	12.0	2022.0
2022-12-06 23:00:00	4.554861	93.283785	6.0	12.0	2022.0


```
[ ]: df.shape
```

```
[ ]: (8893, 10)
```

```
[ ]: plt.figure(dpi=105,figsize=(14,4))
df["P2"].iloc[7000:8893].plot(legend=True)
# df["P2"].iloc[1300000:1316748].plot(legend=True)
# plt.legend(['Training set (40000 Hours)', 'Test set'])
plt.title('PM2.5 Levels')
plt.ylabel("PM2.5")
plt.show()
```



```
[ ]: plt.figure(dpi=105,figsize=(14,4))
df["P1"].iloc[7000:8893].plot(legend=True)
# df["P1"].iloc[1300000:1316748].plot(legend=True)
plt.title('PM10 Levels')
plt.ylabel("PM10")
plt.show()
```



Above we see the plot of PM2.5 and PM10. Looks very similar but still slightly different.

```
[ ]: fig = px.histogram(df, x=df["P1"])

fig.update_layout(
    title={
        'text': "Distribution of PM10 values registered by sensors in various_
↪location in Eindhoven",
        'y':0.95,
        'x':0.5,
        'xanchor': 'center',
        'yanchor': 'top'})
fig.update_yaxes(showgrid=False) # turning off the grid
fig.show()


# -----

fig2 = px.histogram(df, x=df["P2"])


fig2.update_layout(
    title={
        'text': "Distribution of PM2.5 values registered by sensors in various_
↪location in Eindhoven",
        'y':0.95,
        'x':0.5,
        'xanchor': 'center',
        'yanchor': 'top'})
fig2.update_yaxes(showgrid=False) # turning off the grid
fig2.show()
```

Below we will use a heatmap to have a overview of the correlation between the fetaures. The darker the color, the more correlation.


```
[ ]: plt.figure(figsize=(15,15))
cor = df.corr(method='pearson')
sns.heatmap(cor, annot=True, cmap=plt.cm.Reds)
plt.show()
```



```
[ ]: plt.figure(figsize=(15,15))
cor = df.corr(method='spearman')
sns.heatmap(cor, annot=True, cmap=plt.cm.Greens)
plt.show()
```

```
[ ]: plt.figure(figsize=(15,15))
cor = df.corr(method='kendall')
sns.heatmap(cor, annot=True, cmap=plt.cm.Blues)
plt.show()
```



The PM values are heavily correlated with each other, as expected.

Splitting datasets

```
[ ]: df.columns
[ ]: Index(['hour', 'sensor_id', 'P1', 'P2', 'pressure', 'temperature', 'humidity',
          'day', 'month', 'year'],
          dtype='object')
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8893 entries, 2021-12-01 11:00:00 to 2022-12-06 23:00:00
Freq: H
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   hour             8330 non-null   float64
1   sensor_id        8330 non-null   float64
2   P1               8330 non-null   float64
3   P2               8330 non-null   float64
4   pressure         8309 non-null   float64
5   temperature      8309 non-null   float64
6   humidity         8309 non-null   float64
7   day              8330 non-null   float64
8   month            8330 non-null   float64
9   year             8330 non-null   float64
dtypes: float64(10)
memory usage: 764.2 KB
```

```
[ ]: df.shape
```

```
[ ]: (8893, 10)
```

Scale the dataset for LSTM

```
[ ]: min_max_scaler = MinMaxScaler(feature_range=(0,1))
for col in df.columns:
    if col == 'P2' and col == 'P1':
        continue
    else:
        df[col] = min_max_scaler.fit_transform(df[[col]])
```

```
[ ]: # convert multiple records into 1 record having history of last n hours of data
      ↪ as attributes
def transform_data_many_to_one (data, columns, time_steps=1):
    n_vars = data.shape[1]
    dataset = pd.DataFrame(data)
    cols, names = list(), list()
    for i in range(time_steps, 0, -1):
        cols.append(dataset.shift(i))
        names += [('{}(t-{})'.format(columns[j], i)) for j in range(n_vars)]
    cols.append(dataset.shift(-0))
    names += [('{}(t)'.format(columns[j])) for j in range(n_vars)]
    new_df = pd.concat(cols, axis=1)
    new_df.columns = names
    new_df.dropna(inplace=True)
```

```
return new_df
```

```
[ ]: df.head()
```

```
[ ]:          hour  sensor_id      P1      P2  pressure \
timestamp
2021-12-01 11:00:00  0.478261  0.002918  0.100451  0.021469      NaN
2021-12-01 12:00:00  0.521739  0.002918  0.111201  0.019940      NaN
2021-12-01 13:00:00  0.565217  0.002918  0.130084  0.024811      NaN
2021-12-01 14:00:00  0.608696  0.002918  0.245531  0.059394      NaN
2021-12-01 15:00:00  0.652174  0.002918  0.305356  0.078062      NaN

          temperature  humidity  day  month  year
timestamp
2021-12-01 11:00:00      NaN      NaN  0.0    1.0    0.0
2021-12-01 12:00:00      NaN      NaN  0.0    1.0    0.0
2021-12-01 13:00:00      NaN      NaN  0.0    1.0    0.0
2021-12-01 14:00:00      NaN      NaN  0.0    1.0    0.0
2021-12-01 15:00:00      NaN      NaN  0.0    1.0    0.0
```

```
[ ]: # remove unused columns.
df.drop(['sensor_id', 'day', 'month', 'year', 'hour', 'pressure',
        ↪ 'temperature', 'humidity'], axis=1, inplace=True)
```

```
[ ]: df.head()
```

```
[ ]:          P1      P2
timestamp
2021-12-01 11:00:00  0.100451  0.021469
2021-12-01 12:00:00  0.111201  0.019940
2021-12-01 13:00:00  0.130084  0.024811
2021-12-01 14:00:00  0.245531  0.059394
2021-12-01 15:00:00  0.305356  0.078062
```

```
[ ]: values = df.values
values = values.astype('float32')
n_hours = 1
transformed_df = transform_data_many_to_one (values, df.columns, n_hours)
transformed_df.drop(['P1(t)', 'P2(t)'], axis=1, inplace=True)
transformed_df.reset_index (drop=True, inplace=True)
```

```
[ ]: transformed_df.shape
```

```
[ ]: (8237, 2)
```

```
[ ]: transformed_df.head()
```

```
[ ]:      P1(t-1)   P2(t-1)
0  0.100451  0.021469
1  0.111201  0.019940
2  0.130084  0.024811
3  0.245531  0.059394
4  0.305356  0.078062
```

```
[ ]: values = transformed_df.values
n_features = transformed_df.shape[1]
n_train_hours = int(len(transformed_df)*0.8)
n_attributes = n_hours * n_features

train = values[:n_train_hours, :]
# validate = values[n_train_hours:, :] #kept 1 year data for validation
test = values[n_train_hours:, :] # 1 year data for test

train_X, train_y = train[:, :n_attributes], train[:, -1]
test_X, test_y = test[:, :n_attributes], test[:, -1]
# validate_X, validate_y = validate[:, :n_attributes], validate[:, -1]

train_X = train_X.reshape((train_X.shape[0], n_hours, n_features))
# validate_X = validate_X.reshape((validate_X.shape[0], n_hours, n_features))
test_X = test_X.reshape((test_X.shape[0], n_hours, n_features))

# Print data shape
print('number of attributes: ', n_attributes)
print('number of features: ', n_features)
print('number of hours;/window/lookback: ', n_hours)
print('train_X.shape: ', train_X.shape)
print('train_y.shape: ', train_y.shape)
print('test_X.shape: ', test_X.shape)
print('test_y.shape: ', test_y.shape)
# print('validate_X.shape: ', validate_X.shape)
# print('validate_y.shape: ', validate_y.shape)
```

```
number of attributes: 2
number of features: 2
number of hours;/window/lookback: 1
train_X.shape: (6589, 1, 2)
train_y.shape: (6589,)
test_X.shape: (1648, 1, 2)
test_y.shape: (1648,)
```

Model

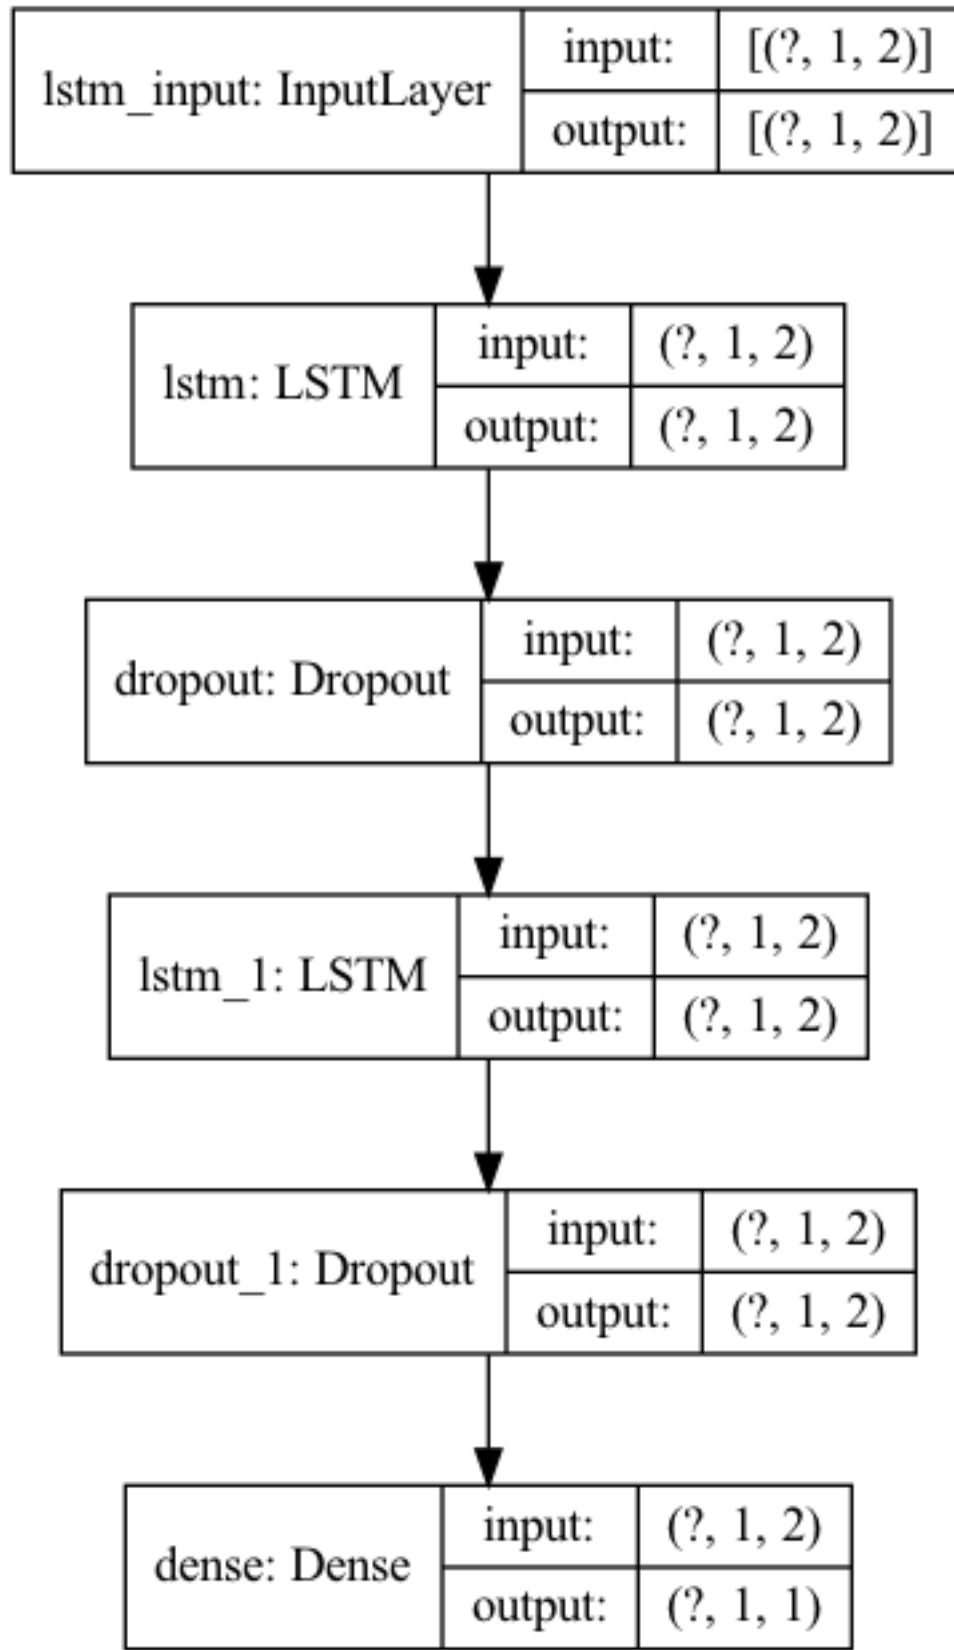
```
[ ]: model = Sequential()
model.add(LSTM(2, return_sequences = True, input_shape=(n_hours, n_features)))
model.add(Dropout(0.2))
```

```
model.add(LSTM(2, return_sequences = True))
model.add(Dropout(0.2))
model.add(Dense(1))
model.compile(optimizer = 'adam', loss='mse')

tf.keras.utils.plot_model(model=model, show_shapes=True)
```

```
2023-01-11 15:46:18.153146: I tensorflow/core/platform/cpu_feature_guard.cc:145]
This TensorFlow binary is optimized with Intel(R) MKL-DNN to use the following
CPU instructions in performance critical operations: SSE4.1 SSE4.2
To enable them in non-MKL-DNN operations, rebuild TensorFlow with the
appropriate compiler flags.
2023-01-11 15:46:18.153636: I
tensorflow/core/common_runtime/process_util.cc:115] Creating new thread pool
with default inter op setting: 8. Tune using inter_op_parallelism_threads for
best performance.
```

```
[ ]:
```



Train model and plot evaluation results

```
[ ]: def model_train_evaluation(X, y, model, model_name):
    #Model run

    # early_stop = EarlyStopping(monitor='val_loss', mode='min', verbose=1,
    ↪patience=3)
    # lr_monitor = tf.keras.callbacks.ReduceLROnPlateau(monitor="val_loss",
    ↪patience=3, factor=0.5, cooldown=1)
    history = model.fit(train_X, train_y, epochs=5, batch_size=16,
    ↪validation_split=0.2, verbose=1, shuffle=False)
    history_frame = pd.DataFrame(history.history)
    print('\n \n')

    # Model Evaluation metrics
    ypred = model.predict(X).flatten()
    print("LSTM Model Evaluation Report: ")
    print('Mean Absolute Error(MAE) of', model_name,':', mean_absolute_error(y,
    ↪ypred))
    print('Mean Absolute Percentage Error (MAPE) of', model_name,':',
    ↪mean_absolute_percentage_error(y, ypred))
    print('Mean Squared Error(MSE) of', model_name,':', mean_squared_error(y,
    ↪ypred))
    print('Root Mean Squared Error (RMSE) of', model_name,':',
    ↪mean_squared_error(y, ypred, squared = False))
    # print('Explained Variance Score (EVS) of', model_name,':',
    ↪explained_variance_score(y, ypred))
    print('R2 Score of', model_name,':', (r2_score(y, ypred)).round(2))
    print('\n \n')

    # model performance plot
    plt.figure(figsize=(20,5),dpi=100)
    plt.plot(history.history['loss'], label='training loss')
    plt.plot(history.history['val_loss'], label='validation loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(loc='best')
    plt.title(model_name + ' Performance Loss')
    plt.show()
    print('\n \n')

    # Actual vs Predicted Plot
    f, ax = plt.subplots(figsize=(12,6),dpi=100);
    plt.scatter(y, ypred, label="Actual vs Predicted")
    # predictions
```



```

plt.xlabel('PM2.5')
plt.ylabel('PM2.5')
plt.title('LSTM: Expection vs Prediction')
plt.plot(y,y,'r', label="Expected Prediction")
plt.legend()
f.text(0.95, 0.06, 'author: AIRQ',
      fontsize=10, color='green',
      ha='left', va='bottom', alpha=0.5);
print('\n \n')

# Plot test data vs prediction
plt.figure(dpi=100, figsize = (11, 7))
range_future = len(ypred)
plt.plot(np.arange(range_future), np.array(ypred),label='Predicted values')
plt.plot(np.arange(range_future), np.array(y), label='Actual values')
plt.title('Prediction vs Actual for ' + model_name)
plt.legend(['Predicted', 'Actual'], loc='upper right')
plt.ylabel('Values')
print('\n \n')

model_train_evaluation(test_X, test_y, model, 'LSTM Model')

```

Train on 5271 samples, validate on 1318 samples

Epoch 1/5

2023-01-11 15:46:21.686615: W

tensorflow/core/grappler/optimizers/implementation_selector.cc:310] Skipping optimization due to error while loading function libraries: Invalid argument: Functions '__inference__backward_standard_lstm_4880_5365_specialized_for_StatefulPartitionedCall_at__inference_distributed_function_6006' and '__inference__backward_cudnn_lstm_with_fallback_4525_4707' both implement 'lstm_6c52b04a-11ea-4ede-8c99-4b38f62d430d' but their signatures do not match.

5264/5271 [=====>.] - ETA: 0s - loss: 0.0878

2023-01-11 15:46:26.873905: W

tensorflow/core/grappler/optimizers/implementation_selector.cc:310] Skipping optimization due to error while loading function libraries: Invalid argument: Functions '__inference_cudnn_lstm_with_fallback_6967' and '__inference_standard_lstm_6856_specialized_for_sequential_lstm_StatefulPartitionedCall_at__inference_distributed_function_7680' both implement 'lstm_9bc4594b-4523-4918-b321-b85eeb838da8' but their signatures do not match.

5271/5271 [=====] - 9s 2ms/sample - loss: 0.0878 - val_loss: 0.0344

Epoch 2/5

5271/5271 [=====] - 2s 446us/sample - loss: 0.0493 - val_loss: 0.0263


Epoch 3/5



```
5271/5271 [=====] - 2s 447us/sample - loss: 0.0357 -  
val_loss: 0.0156  
Epoch 4/5  
5271/5271 [=====] - 2s 448us/sample - loss: 0.0257 -  
val_loss: 0.0081  
Epoch 5/5  
5271/5271 [=====] - 2s 442us/sample - loss: 0.0189 -  
val_loss: 0.0050
```

2023-01-11 15:46:37.664784: W
tensorflow/core/grappler/optimizers/implementation_selector.cc:310] Skipping
optimization due to error while loading function libraries: Invalid argument:
Functions '__inference_standard_lstm_11396' and '__inference_standard_lstm_11396
_specialized_for_sequential_lstm_StatefulPartitionedCall_at___inference_distribu
ted_function_12193' both implement 'lstm_7aaa2bc5-9f60-49d0-af4f-fd7a2757ca2d'
but their signatures do not match.

LSTM Model Evaluation Report:

Mean Absolute Error(MAE) of LSTM Model : 0.060889676
Mean Absolute Percentage Error (MAPE) of LSTM Model : 0.51925695
Mean Squared Error(MSE) of LSTM Model : 0.00510778
Root Mean Squared Error (RMSE) of LSTM Model : 0.07146873
R2 Score of LSTM Model : 0.92





Save model

```
[ ]: model.save("lstm_model_v5.h5")
```

Load model and check values. This block of code can also be used on a validation set to check the predicted value against the actual values.

```
[ ]: loaded_model = tf.keras.models.load_model('lstm_model_v3.h5')
train_predictions = loaded_model.predict(train_X).flatten()
train_results = pd.DataFrame(data={'Train Predictions': train_predictions,
    ↪ 'Actual': train_y})
train_results.head()
```

```
[ ]: 
```

	Train Predictions	Actual
0	0.810553	0.899267
1	0.801528	0.854910
2	0.799165	0.835224
3	0.791139	0.792853
4	0.775569	0.746143